

Aly Sultan

978-325-1925 | Boston, MA

sultan.a@northeastern.edu | [linkedin.com/in/aly-sultan/](https://www.linkedin.com/in/aly-sultan/) | github.com/asultan123

EDUCATION

Northeastern University

Ph.D. Computer Engineering

Boston, MA

2019 – Expected 2025

American University in Cairo

B.S. Electronics & Communication Engineering

Cairo, Egypt

2014 – 2019

EXPERIENCE

Graduate Research Assistant

2020 – Present

Embedded Systems Lab, Northeastern University

Boston, MA

- Developed a SystemC model of a novel Hybrid general matrix multiplication and convolution accelerator (HERO) for accelerating inference in deep neural networks
- Developed a Template Optimization tool (TEMPO) in python that optimizes HERO templates based on available compute resources and a target library of neural networks
- Estimated HERO's latency and energy consumption using a Python-SystemC based simulation environment (HERO-Sim)
- Created a SystemC model of a novel programmable memory primitive called Self Addressable Memory (SAM) used in statically scheduling dataflows in neural network accelerator architectures
- Implemented a SAM program compiler that generates data movement programs from convolution layer descriptions to orchestrate on-chip data movement between SAMS in HERO
- Supervised and mentored team of three undergraduate students in:
 - * Implementing SAMs in Verilog
 - * Integrating HERO into Xilinx's SystemC+QEMU simulation environment

PROJECTS

Integer Linear Program based scheduler for multi-core processors | *Python, Pyomo, Gurobi*

2021

- Implemented an ILP-based scheduling model to schedule applications statically on a multicore processor
- Model was generated using python and the pyomo library and optimized with the Gurobi solver
- Successfully generated optimal schedules for a variety of application sizes and core count configurations

Cache Optimization for CNN Inference | *C, Darknet, Intel Pin*

2019

- Used Intel Pin and the Darknet framework to determine the effect of cache configurations on CNN inference
- Explored effects of cache hierarchy levels, sizing and replacement policy on data movement to and from DRAM during CNN inference

Darknet Convolution Inference accelerator | *C, Darknet, Vivado HLS*

2019

- Developed a General Matrix Multiplication (GEMM) accelerator using Vivado HLS
- Synthesized accelerator on a Zynq-7020 FPGA
- Integrated accelerator into Darknet framework
- Accelerated inference time of tiny darknet network by a factor of 2X over CPU Baseline

PUBLICATIONS

A Compact Low-Power Mitchell-Based Error Tolerant Multiplier | *Verilog, Matlab*

2018

- Developed a novel approximate multiplier architecture in Verilog and evaluated numerical accuracy in JPEG compression using MATLAB
- Design improved power-delay-product by 1.9X with only 20% reduction in peak signal-to-noise ratio in JPEG compression compared to a Xilinx Zynq-7020 DSP
- Published and presented findings at NGCAS, Malta

TECHNICAL SKILLS

Languages: Python, C/C++, SystemC, Matlab, Verilog

Framework: Intel Pin, Darknet

Developer Tools: Git, Docker, VS Code, Vivado, Vivado HLS, QEMU, Gurobi

Libraries: PyTorch, Numpy, Matplotlib, Pyomo